


# Databases

Software Development & Python  
Albina Rukhadze, 2022




**girafe**  
**ai**





# SQL basics



# SQL Data Types



## Integer Types

- **smallint** – signed two-byte integer
- **integer** – signed four-byte integer
- **bigint** – signed eight-byte integer

## Arbitrary Precision Numbers

- **numeric(p, s) (= decimal(p, s))**  
– exact numeric of selectable precision

## Floating-Point Types

- **real (= float4)** – single precision floating-point number (4 bytes)
- **double precision (= float8)** – double precision floating-point number (8 bytes)



# SQL Data Types



## Character Types

- **character varying(n), varchar(n)** – variable-length with limit
- **character(n), char(n)** – fixed-length, blank padded
- **text** – variable unlimited length

## Date/Time Types

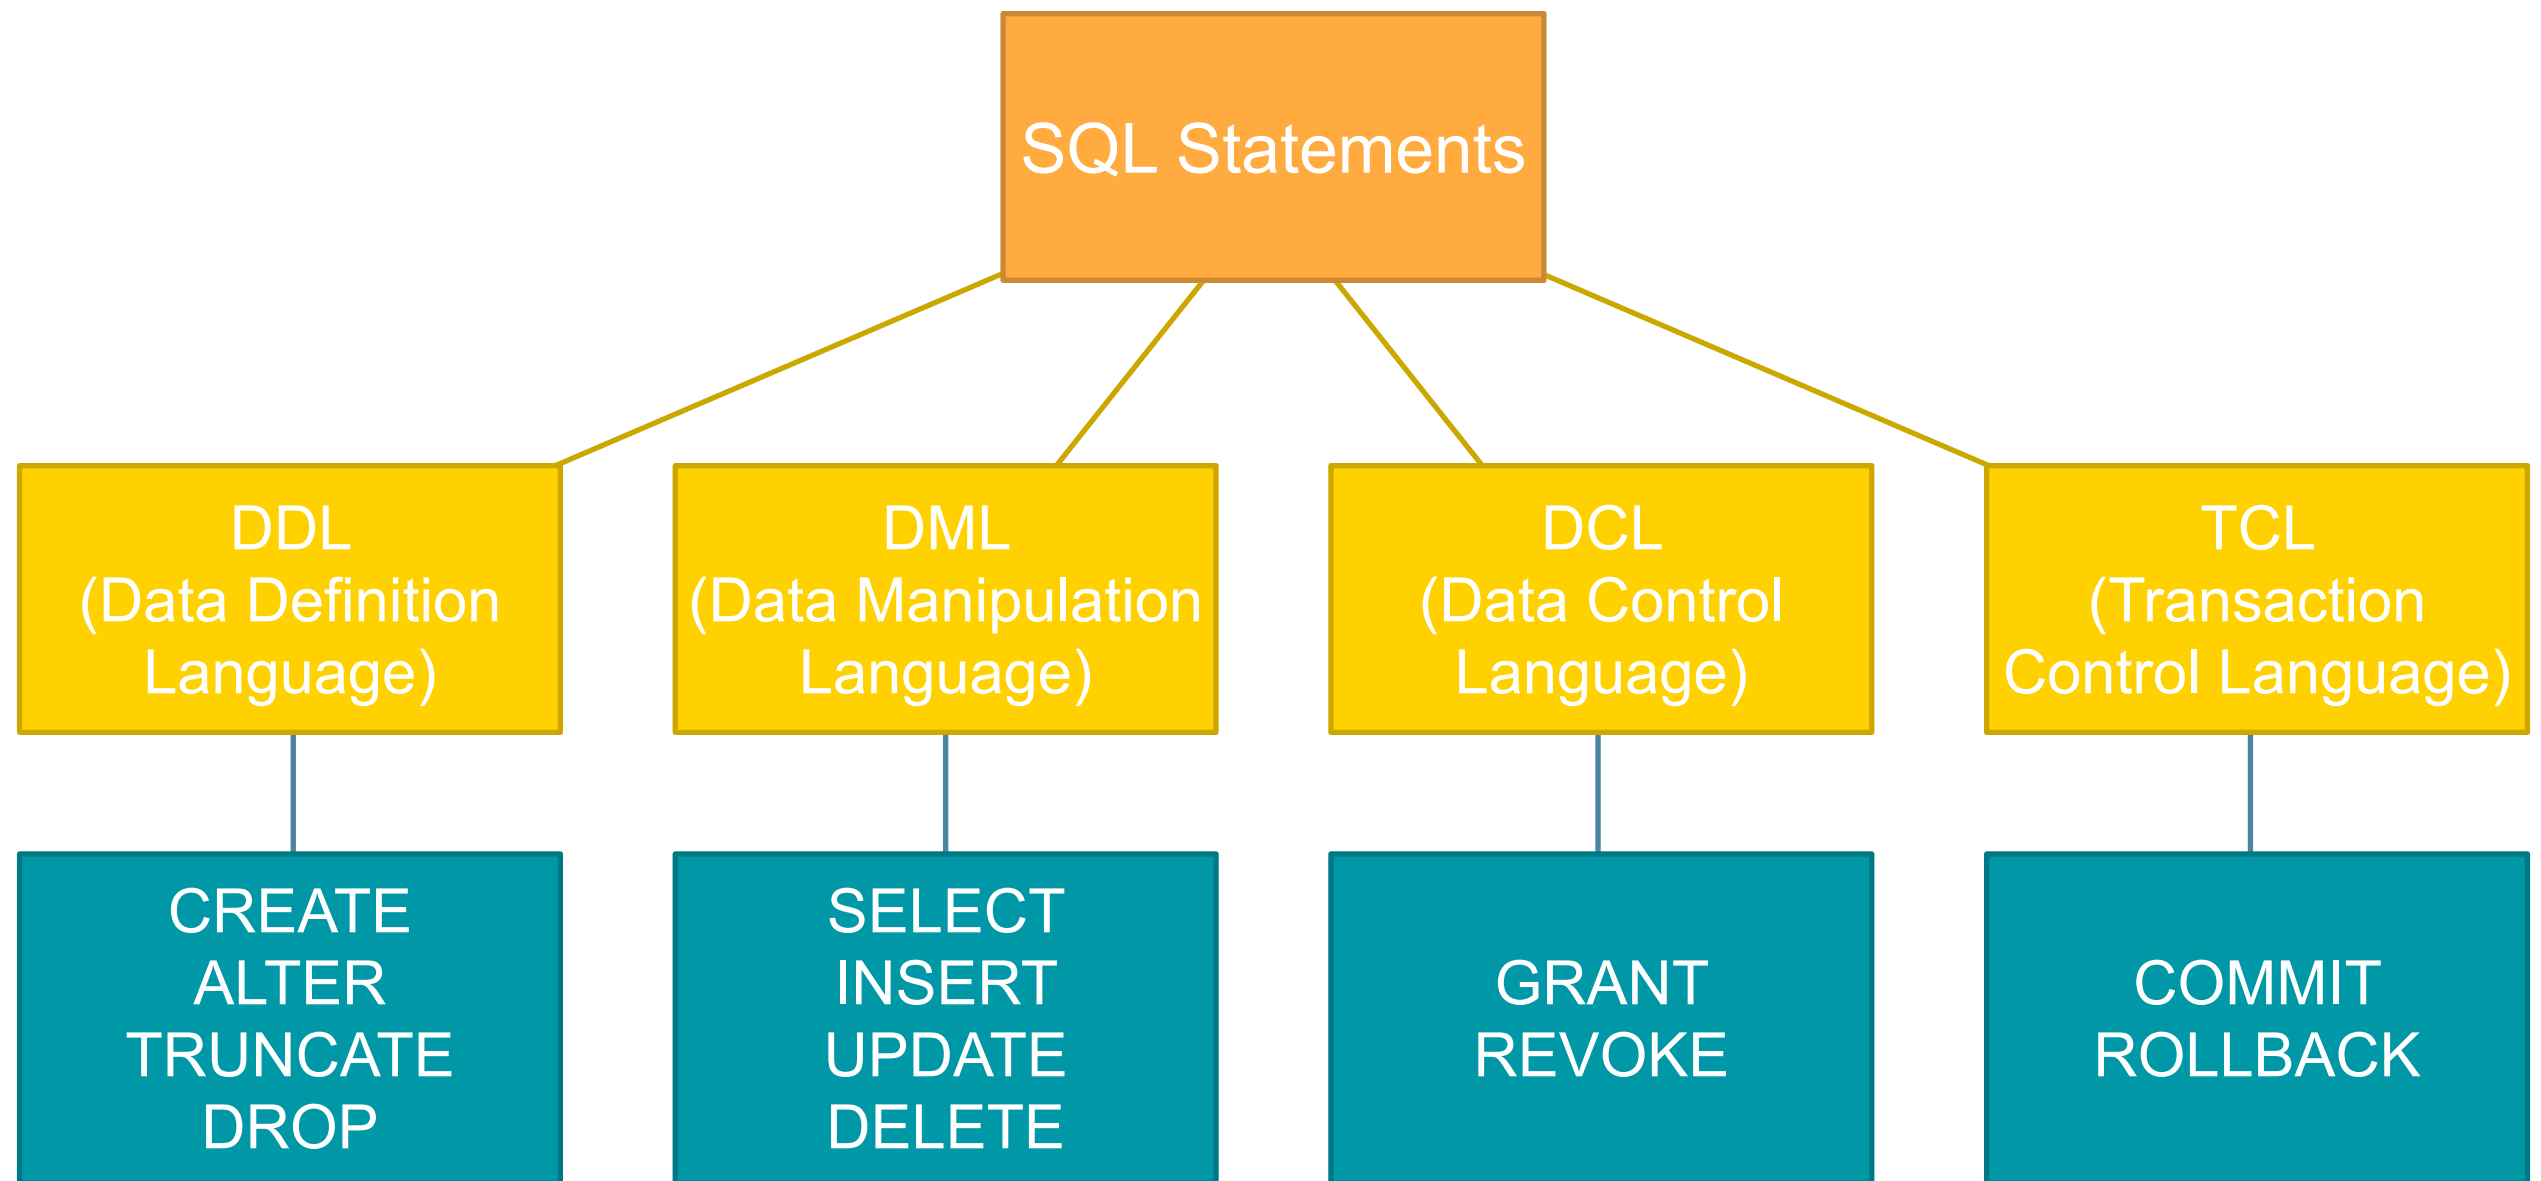
- **date** – date (no time of day)
- **time** – time of day (no date)
- **timestamp** – both date and time
- **time with timezone** – time of day (no date), with time zone
- **timestamp with timezone** – both date and time, with time zone
- **interval** – time interval

## Boolean Types

- **boolean** – state of true or false



# SQL Operators



# Data Definition Language

Database objects creation, modification  
and deleting

---

**girafe**  
**ai**

# CREATE



– define a new table

```
CREATE [TEMPORARY] TABLE [IF NOT EXISTS] table_name(  
    col_name_1    datatype_1,  
    col_name_2    datatype_2,  
    ...  
    col_name_N    datatype_N  
);
```

```
CREATE TABLE films (  
    code          char(5) PRIMARY KEY,  
    title         varchar(40) NOT NULL,  
    did          integer NOT NULL,  
    date_prod     date,  
    kind          varchar(10),  
    len          interval hour to minute  
);
```



# ALTER



– change the definition of a table

```
ALTER TABLE table_name ADD column_name datatype;  
ALTER TABLE table_name DROP column_name;  
ALTER TABLE table_name RENAME column_name TO new_name;  
ALTER TABLE table_name ALTER column_name TYPE datatype;  
...
```

```
ALTER TABLE distributors ADD COLUMN address varchar(30);
```

```
ALTER TABLE transactions  
  ADD COLUMN status varchar(30) DEFAULT 'old',  
  ALTER COLUMN status SET default 'current';
```





# DROP



– remove a table

```
DROP TABLE [IF EXISTS] table_name;
```

```
DROP TABLE films, distributors;
```



# TRUNCATE



– empty a table or set of tables (remove all rows from a set of tables)

```
TRUNCATE TABLE table_name;
```

```
TRUNCATE bigtable, fattable;
```



# Data Manipulation Language

Inserting, deleting and modifying data  
in a database

---

**girafe**  
**ai**

# SELECT



– retrieve rows from a table or view

```
SELECT [DISTINCT] select_item_comma_list    -- columns output list
FROM table_reference_comma_list             -- table list
-- filter conditions, AND/OR/NOT may be used
[WHERE conditional_expression]
[GROUP BY column_name_comma_list]           -- grouping condition
-- filter conditions after grouping
[HAVING conditional_expression]
[ORDER BY order_item_comma_list];           -- output sorting columns
```

```
SELECT *
FROM films
ORDER BY title;
```

## Query execution order

**FROM** → **WHERE** → **GROUP BY** → **HAVING** → **SELECT** → **ORDER BY**



# INSERT



– create new rows in a table

```
INSERT INTO table_name [(comma_separated_column_names)]  
    VALUES (comma_separated_values);
```

```
INSERT INTO films VALUES  
    ('UA502', 'Bananas', 105, DEFAULT,  
    'Comedy', '82 minutes');  
INSERT INTO films (code, title, did, date_prod, kind)  
    VALUES ('T_601', 'Yojimbo', 106, DEFAULT, 'Drama');
```



# UPDATE



– update rows of a table

```
UPDATE table_name  
    SET update_assignment_comma_list  
WHERE conditional_expression;
```

```
UPDATE films  
    SET kind = 'Dramatic'  
WHERE kind = 'Drama';
```



# DELETE



– delete rows of a table

**DELETE**

**FROM** `table_name`

[**WHERE** `conditional_expression`];

**DELETE**


**FROM** `films`

**WHERE** `kind <> 'Musical'`;





# Profound SQL





# Aggregate Functions

Calculating statistics on a set of values

---

**girafe**  
**ai**

# Standard Aggregate Functions



– when grouping, the **SELECT** block may contain either attributes by which grouping occurs, or attributes that are input to aggregating functions.

- **count()** – the number of records with a known value, **count(DISTINCT field\_nm)** may be used to count the number of unique values;
- **max()** – the largest of all selected field values;
- **min()** – the least of all selected field values;
- **sum()** – sum of all selected field values;
- **avg()** – average of all selected field values.

**Note:** **NULL**-values are ignored



# Joining Tables

Merge tables on a given condition

---

**girafe**  
**ai**

# Join Operators



Join operators are divided into 3 groups:


- **CROSS JOIN** – Cartesian product of 2 tables;
- **INNER JOIN** – joining 2 tables by condition. Query result will include only records that satisfy the connection condition;
- **OUTER JOIN** – joining 2 tables by condition. Query result may include records that do not satisfy the connection condition:
  - **LEFT (OUTER) JOIN** – result contains all rows of the "left" table;
  - **RIGHT (OUTER) JOIN** – result contains all rows of the "right" table;
  - **FULL (OUTER) JOIN** – result contains all rows of both tables.




# Join Operators




- **CROSS JOIN**
- **INNER JOIN**
- **OUTER JOIN**
  - **LEFT (OUTER) JOIN**
  - **RIGHT (OUTER) JOIN**
  - **FULL (OUTER) JOIN**




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```



# Join Operators



```
SELECT *  
FROM weather INNER JOIN cities  
      ON city = name;
```

city	temp_lo	temp_hi	prcp	date	name	location
San Francisco	46	50	0.25	1994-11-27	San Francisco	(-194, 53)
San Francisco	43	57	0	1994-11-29	San Francisco	(-194, 53)

(2 rows)



# Join Operators



```
SELECT *  
FROM weather LEFT OUTER JOIN cities  
      ON weather.city = cities.name;
```

city	temp_lo	temp_hi	prcp	date	name	location
Hayward	37	54		1994-11-29		
San Francisco	46	50	0.25	1994-11-27	San Francisco	(-194, 53)
San Francisco	43	57	0	1994-11-29	San Francisco	(-194, 53)

(3 rows)



# Join Operators



```
SELECT *  
FROM weather RIGHT OUTER JOIN cities  
      ON weather.city = cities.name;
```

city	temp_lo	temp_hi	prcp	date	name	location
					Tokyo	(139, 35)
San Francisco	46	50	0.25	1994-11-27	San Francisco	(-194, 53)
San Francisco	43	57	0	1994-11-29	San Francisco	(-194, 53)


(3 rows)








# Transactions



# Transaction



– an object that groups a sequence of operations that must be performed as a whole. Ensures the transition of the database from one consistent state to another.



Transactions ensure the integrity of the database under the conditions:

- Parallel data processing;
- Physical disk failures;
- Emergency power failure.



# ACID

A set of properties of database transactions intended to guarantee data validity despite errors, power failures, and other mishaps

---

**girafe**  
**ai**

# Atomicity (ACID)



- A transaction must be an atomic (indivisible) unit of work
- Either all operations included in the transaction must be performed, or none of them
- Therefore, if it is impossible to perform all operations, all the changes made must be cancelled



# Consistency (ACID)



- Upon completion of the transaction, all data should remain in a consistent state
- When executing a transaction, all the rules of a relational DBMS must be followed:
  - Checks for compliance with restrictions (domains, uniqueness indexes, foreign keys, checks, rules, etc.)
  - Updating indexes;
  - Executing triggers
  - etc.



# Isolation (ACID)



- Data changes performed within a transaction must be isolated from all changes performed in other transactions until the transaction is committed
- There are different levels of isolation – to achieve a compromise between the degree of parallelization of work with the database and the rigor of the principle of consistency:
  - The higher the isolation level, the higher the degree of data consistency
  - The higher the isolation level, the lower the degree of parallelization and the lower the degree of data availability



# Durability (ACID)



- If a transaction has been completed, its result should be saved in the system, despite the system failure
- If the transaction has not been completed, its result must be completely canceled following a system failure



# Transaction Control Language

Transactions managing

---

**girafe**  
**ai**



# BEGIN



– set a transaction block

```
BEGIN TRANSACTION transaction_mode [, ...]  
    ISOLATION LEVEL { SERIALIZABLE | REPEATABLE READ |  
                      READ COMMITTED | READ UNCOMMITTED }  
    READ WRITE | READ ONLY  
    [NOT] DEFERRABLE
```



# COMMIT



– commit the current transaction

**COMMIT [ WORK | TRANSACTION ] [ AND [ NO ] CHAIN ]**



# ROLLBACK



– abort the current transaction

**ROLLBACK [ WORK | TRANSACTION ] [ AND [ NO ] CHAIN ]**



# SAVEPOINT



- define a new savepoint within the current transaction

**SAVEPOINT** savepoint\_name



# Transaction Operators



```
BEGIN;  
    INSERT INTO table_name VALUES (1);  
    SAVEPOINT my_savepoint;  
    INSERT INTO table_name VALUES (2);  
    ROLLBACK TO SAVEPOINT my_savepoint;  
    INSERT INTO table_name VALUES (3);  
COMMIT;
```




# Transaction Isolation

Classification of problems and isolation levels

---

**girafe**  
**ai**

# Problem Classification



# Dirty read



- a transaction reads data written by a concurrent uncommitted transaction.

Transaction 1	Transaction 2
<b>UPDATE</b> table_1 <b>SET</b> attr_2 = attr_2 + 20 <b>WHERE</b> attr_1 = 1;	
	<b>SELECT</b> attr_2 <b>FROM</b> table_1 <b>WHERE</b> attr_1 = 1;
<b>ROLLBACK WORK;</b>	





# Nonrepeatable read



- a transaction re-reads data it has previously read and finds that data has been modified by another transaction (that committed since the initial read).

Transaction 1	Transaction 2
	<b>SELECT</b> attr_2 <b>FROM</b> table_1 <b>WHERE</b> attr_1 = 1;
<b>UPDATE</b> table_1 <b>SET</b> attr_2 = attr_2 + 20 <b>WHERE</b> attr_1 = 1;	
<b>COMMIT;</b>	
	<b>SELECT</b> attr_2 <b>FROM</b> table_1 <b>WHERE</b> attr_1 = 1;



# Phantom read



- a transaction re-executes a query returning a set of rows that satisfy a search condition and finds that the set of rows satisfying the condition has changed due to another recently-committed transaction.

Transaction 1	Transaction 2
	<b>SELECT</b> <i>sum(attr_2)</i> <b>FROM</b> table_1;
<b>INSERT INTO</b> table_1 (attr_1, attr_2) <b>VALUES</b> (15, 20);	
<b>COMMIT;</b>	
	<b>SELECT</b> <i>sum(attr_2)</i> <b>FROM</b> table_1;



# Serialization anomaly




- the result of successfully committing a group of transactions is inconsistent with all possible orderings of running those transactions one at a time.

Transaction 1	Transaction 2
<pre>UPDATE table_1   SET attr_2 = attr_2 + 20 WHERE attr_1 = 1;</pre>	<pre>UPDATE table_1   SET attr_2 = attr_2 + 25 WHERE attr_1 = 1;</pre>



# Isolation Level



# Isolation Levels




Isolation Level	Dirty Read	Nonrepeatable Read	Phantom Read	Serialization Anomaly
0. Read Uncommitted	possible	possible	possible	possible
1. Read Committed	-	possible	possible	possible
2. Repeatable Read	-	-	possible	possible
3. Serializable	-	-	-	-





# Databases in Python

DB types, ORM libraries



# Basic Concepts



**Python DB API** – the standard of interfaces for packages working with the database. A set of rules that individual modules implementing work with specific databases must comply with.

Database	Interface
SQLite	<code>sqlite3</code>
PostgreSQL	<code>psycopg2</code>
MySQL	<code>PyMySQL</code>
Oracle	<code>cx_Oracle</code>



# Basic Concepts



## **Benefits of ORM usage:**

- no need to write queries in pure SQL
- developer may abstract from thoughts about the internal structure of the database

## **Possible difficulties:**

- ORM may work much slower than direct interaction with the database
- Writing queries in ORM syntax can be technically much more difficult





# ORM (Object Relational Mapping)



**Python DB API** – the standard of interfaces for packages working with the database. A set of rules that individual modules implementing work with specific databases must comply with.

Database	Interface
SQLite	<code>sqlite3</code>
PostgreSQL	<code>psycopg2</code>
MySQL	<code>PyMySQL</code>
Oracle	<code>cx_Oracle</code>



# SQLite3



## Main features:

- Lightweight solution to use a database
- All database is stored in one file
- Suitable for mobile apps
- Non-scalable solution



# Peewee



## Main features:

- Built-in support for SQLite, MySQL and PostgreSQL
- Lightweight, flexible and expressive
- Low entry threshold



# SQLAlchemy



## **Advantages:**

- SQL injection protection: query parameters are escaped
- Scalability and flexibility
- Ability to use both pure SQL and ORM syntax
- Large-scale solution



